IDS 690: Unifying Data Science Interpreting Indicator Variables

Derek Wales and Prajwal Vijendra

In [1]:
```
import pandas as pd
import numpy as np
import statsmodels.api as sm
import statsmodels.formula.api as smf
from statsmodels.formula.api import ols
import matplotlib.pyplot as plt
```

In [2]:
```
# Load in data
auto_df = pd.read_stata('automobile_dataset.dta')
auto_df.head()
```

Out[2]:

	make	price	mpg	rep78	headroom	trunk	weight	length	turn	displacement	gear_ratio
0	AMC Concord	4099	22	3.0	2.5	11	2930	186	40	121	3.58
1	AMC Pacer	4749	17	3.0	3.0	11	3350	173	40	258	2.53
2	AMC Spirit	3799	22	NaN	3.0	12	2640	168	35	121	3.08
3	Buick Century	4816	20	3.0	4.5	16	3250	196	40	196	2.93
4	Buick Electra	7827	15	4.0	4.0	20	4080	222	43	350	2.41

Exercise 1

In [3]:
```
# Create new indicator variable
# 1 indicates mpg < 18
auto_df['guzzler'] = np.where(auto_df['mpg']<18,1,0)
```

In [4]:
```
# Linear regression on car price vs. guzzler
smf.ols('price ~ C(guzzler)', data = auto_df).fit().summary()
```

Out[4]:

OLS Regression Results

Dep. Variable:	price	**R-squared:**	0.379
Model:	OLS	**Adj. R-squared:**	0.370
Method:	Least Squares	**F-statistic:**	43.90
Date:	Fri, 21 Feb 2020	**Prob (F-statistic):**	5.38e-09
Time:	13:13:11	**Log-Likelihood:**	-678.10
No. Observations:	74	**AIC:**	1360.
Df Residuals:	72	**BIC:**	1365.
Df Model:	1		
Covariance Type:	nonrobust		

	coef	**std err**	**t**	**P>|t|**	**[0.025**	**0.975]**
Intercept	5143.0893	312.807	16.442	0.000	4519.521	5766.658
C(guzzler)[T.1]	4202.2440	634.243	6.626	0.000	2937.904	5466.584

Omnibus:	37.244	**Durbin-Watson:**	1.348
Prob(Omnibus):	0.000	**Jarque-Bera (JB):**	111.225
Skew:	1.565	**Prob(JB):**	7.04e-25
Kurtosis:	8.126	**Cond. No.**	2.50

Warnings:
[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

Our simple linear regression shows that a gas guzzler adds roughly 4202.24 dollars to the price of a car.

Exercise 2

In [7]:
```
# Scatterplot of price of cars by weight colored by guzzler
plt.scatter(y='price', x='weight', label='Not Guzzler', color = 'blue', data=a
uto_df[auto_df.guzzler==0])
plt.scatter(y='price', x='weight', label='Guzzler', color = 'yellow', data=aut
o_df[auto_df.guzzler==1])
plt.ylabel('Price')
plt.xlabel('Weight')
plt.title('Price by Wieght of Cars')
plt.legend()
plt.show()
```

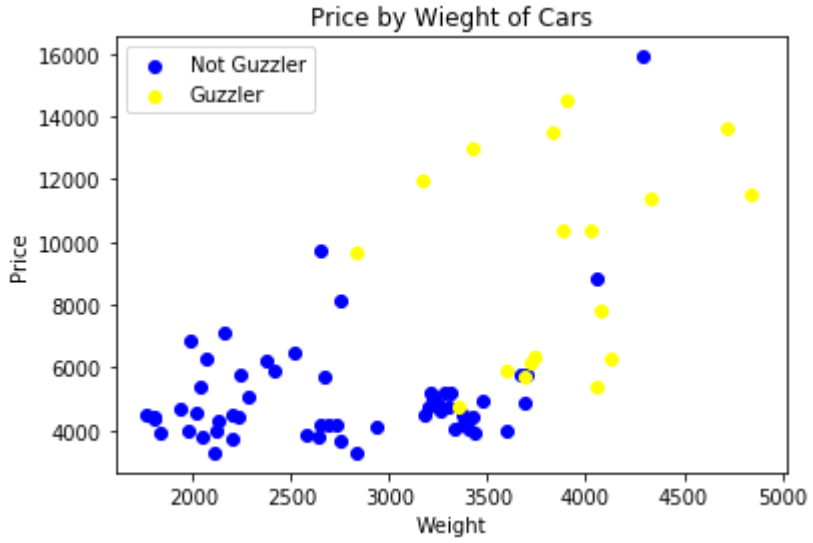

Given our analysis, we should control for weight in our regression. Not controling for weight would overestimate the coefficient for guzzler as most guzzlers tend to be more expensive than non-guzzlers.

Exercise 3

In [6]:
```
# Linear regression on car price vs. guzzler
# Additional controls
smf.ols('price ~ C(guzzler) + weight + C(foreign) + headroom + displacement',
data = auto_df).fit().summary()
```

Out[6]:

OLS Regression Results

Dep. Variable:	price	**R-squared:**	0.596
Model:	OLS	**Adj. R-squared:**	0.566
Method:	Least Squares	**F-statistic:**	20.04
Date:	Wed, 19 Feb 2020	**Prob (F-statistic):**	3.14e-12
Time:	20:18:43	**Log-Likelihood:**	-662.20
No. Observations:	74	**AIC:**	1336.
Df Residuals:	68	**BIC:**	1350.
Df Model:	5		
Covariance Type:	nonrobust		

	coef	std err	t	P>|t|	[0.025	0.975]
Intercept	-782.5353	1612.628	-0.485	0.629	-4000.484	2435.414
C(guzzler)[T.1]	1977.1796	711.055	2.781	0.007	558.291	3396.068
C(foreign)[T.Foreign]	3278.9827	671.826	4.881	0.000	1938.375	4619.591
weight	1.9634	0.702	2.797	0.007	0.563	3.364
headroom	-736.7997	309.009	-2.384	0.020	-1353.418	-120.182
displacement	8.9667	5.819	1.541	0.128	-2.646	20.579

Omnibus:	22.179	**Durbin-Watson:**	1.409
Prob(Omnibus):	0.000	**Jarque-Bera (JB):**	37.284
Skew:	1.118	**Prob(JB):**	8.01e-09
Kurtosis:	5.663	**Cond. No.**	2.36e+04

Warnings:
[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
[2] The condition number is large, 2.36e+04. This might indicate that there are
strong multicollinearity or other numerical problems.

Including additional control variable, our estimate of the guzzler coefficient is reduce to 1,977 from our previous estimate of 4,202. In other words, the additional cost incurred by buying a guzzler is reduced when controling for weight, foreign, headroom, and displacement.

Additionally, controling for guzzler-status, a foreign car costs 3,278 additional dollars on average, each additional lb of a car is associated with a 1.96 dollar increase in price, each unit of headroom is associated with a 736 dollar reduction in cost, and each unit of displacement is associated with a 8.96 dollar increase in cost.

Exercise 4

In [7]:
```
# Create indicator variables for repair record
auto_df = auto_df[pd.notnull(auto_df.rep78)].copy()
auto_df['rep_1'] = np.where(auto_df.rep78 == 1,1,0)
auto_df['rep_2'] = np.where(auto_df.rep78 == 2,1,0)
auto_df['rep_3'] = np.where(auto_df.rep78 == 3,1,0)
auto_df['rep_4'] = np.where(auto_df.rep78 == 4,1,0)
auto_df['rep_5'] = np.where(auto_df.rep78 == 5,1,0)
```

In [8]:
```
# Linear regression on car price vs. guzzler
# Control for repair record
smf.ols('price ~ rep_2 + rep_3 + rep_4 + rep_5 + weight + C(foreign) + headroo
m + displacement', data = auto_df).fit().summary() # C(guzzler)
```

Out[8]:

OLS Regression Results

Dep. Variable:	price	**R-squared:**	0.562
Model:	OLS	**Adj. R-squared:**	0.503
Method:	Least Squares	**F-statistic:**	9.611
Date:	Wed, 19 Feb 2020	**Prob (F-statistic):**	1.87e-08
Time:	20:18:45	**Log-Likelihood:**	-619.34
No. Observations:	69	**AIC:**	1257.
Df Residuals:	60	**BIC:**	1277.
Df Model:	8		
Covariance Type:	nonrobust		

	coef	std err	t	P>|t|	[0.025	0.975]
Intercept	-3674.8333	2181.321	-1.685	0.097	-8038.125	688.458
C(foreign)[T.Foreign]	3565.2581	815.700	4.371	0.000	1933.616	5196.901
rep_2	1292.4864	1717.908	0.752	0.455	-2143.841	4728.814
rep_3	1546.1189	1582.091	0.977	0.332	-1618.534	4710.771
rep_4	1319.9236	1649.062	0.800	0.427	-1978.692	4618.539
rep_5	1917.3066	1732.508	1.107	0.273	-1548.226	5382.839
weight	2.1325	0.890	2.396	0.020	0.352	3.913
headroom	-750.7992	351.685	-2.135	0.037	-1454.274	-47.325
displacement	15.4064	7.517	2.049	0.045	0.369	30.443

Omnibus:	16.791	**Durbin-Watson:**	1.414
Prob(Omnibus):	0.000	**Jarque-Bera (JB):**	19.847
Skew:	1.131	**Prob(JB):**	4.90e-05
Kurtosis:	4.335	**Cond. No.**	4.44e+04

Warnings:
[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
[2] The condition number is large, 4.44e+04. This might indicate that there are
strong multicollinearity or other numerical problems.

According to our analysis, buying a car with an acceptable repair record costs 1,546 dollars more, on average, compared to simialar car (similar weight, headroom, displacement, and place of manufacturing) with a very poor repair record.

Additionally, controling for repair record, a foreign car costs 3,565 additional dollars on average, each additional lb of a car is associated with a 2.13 dollar increase in price, each unit of headroom is associated with a 750 dollar reduction in cost, and each unit of displacement is associated with a 15 dollar increase in cost.

Exercise 5

In [9]:
```
# Linear regression on car price vs. guzzler
# Interaction effects for foreign and guzzler
smf.ols('price ~ weight + C(foreign) + guzzler + C(foreign)*guzzler + headroom
+ displacement', data = auto_df).fit().summary() # C(guzzler)
```

Out[9]:

OLS Regression Results

Dep. Variable:	price	**R-squared:**	0.598
Model:	OLS	**Adj. R-squared:**	0.560
Method:	Least Squares	**F-statistic:**	15.40
Date:	Wed, 19 Feb 2020	**Prob (F-statistic):**	1.02e-10
Time:	20:18:46	**Log-Likelihood:**	-616.32
No. Observations:	69	**AIC:**	1247.
Df Residuals:	62	**BIC:**	1262.
Df Model:	6		
Covariance Type:	nonrobust		

	coef	std err	t	P>|t|	[0.025	0.975]
Intercept	-491.1286	1806.853	-0.272	0.787	-4102.977	3120.720
C(foreign)[T.Foreign]	3064.1019	717.185	4.272	0.000	1630.469	4497.735
weight	1.4155	0.872	1.623	0.110	-0.328	3.159
guzzler	1234.8723	786.967	1.569	0.122	-338.253	2807.997
C(foreign)[T.Foreign]:guzzler	2367.5358	1635.982	1.447	0.153	-902.745	5637.817
headroom	-672.3043	314.885	-2.135	0.037	-1301.751	-42.857
displacement	15.4387	7.492	2.061	0.044	0.463	30.415

Omnibus:	24.422	**Durbin-Watson:**	1.484
Prob(Omnibus):	0.000	**Jarque-Bera (JB):**	40.778
Skew:	1.300	**Prob(JB):**	1.40e-09
Kurtosis:	5.725	**Cond. No.**	2.58e+04

Warnings:
[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
[2] The condition number is large, 2.58e+04. This might indicate that there are
strong multicollinearity or other numerical problems.

According to our analysis, there are additional costs associated with buying a guzzler and buying foreign cars (1,234 and 3,064 respectively). However, there if your car is both a guzzler and foreign, there is an additional 2,367 cost for a car that is both a foreign guzzler on top all previous costs.

Exercise 6

In [10]:
```
foreign = 3064.1019
guzzler = 1234.8723
foreign_x_guzzler = 2367.5358
```

In [11]:
```
# Difference between foreign guzzler and foreign non-guzzler
(foreign + guzzler + foreign_x_guzzler) - (foreign)
```

Out[11]: 3602.4081

Exercise 7

In [12]:
```
# Difference between domestic non-guzzler and foreign non-guzzler
0 - (foreign)
```

Out[12]: -3064.1019

Exercise 8

In [13]:
```
# Linear regression on car price vs. guzzler
# Interaction effects for foreign and guzzler
smf.ols('price ~ C(foreign)*mpg + headroom + weight + displacement', data = au
to_df).fit().summary() # C(guzzler)
```

Out[13]:

OLS Regression Results

Dep. Variable:	price	**R-squared:**	0.586
Model:	OLS	**Adj. R-squared:**	0.546
Method:	Least Squares	**F-statistic:**	14.63
Date:	Wed, 19 Feb 2020	**Prob (F-statistic):**	2.54e-10
Time:	20:18:48	**Log-Likelihood:**	-617.37
No. Observations:	69	**AIC:**	1249.
Df Residuals:	62	**BIC:**	1264.
Df Model:	6		
Covariance Type:	nonrobust		

	coef	std err	t	P>|t|	[0.025	0.975]
Intercept	-1.113e+04	4726.813	-2.355	0.022	-2.06e+04	-1684.462
C(foreign)[T.Foreign]	1.072e+04	3113.992	3.444	0.001	4498.775	1.69e+04
mpg	236.0976	114.808	2.056	0.044	6.599	465.596
C(foreign)[T.Foreign]:mpg	-273.1404	120.012	-2.276	0.026	-513.042	-33.239
headroom	-456.5527	332.445	-1.373	0.175	-1121.101	207.996
weight	2.9201	1.002	2.914	0.005	0.917	4.923
displacement	18.2165	7.228	2.520	0.014	3.769	32.664

Omnibus:	21.421	**Durbin-Watson:**	1.499
Prob(Omnibus):	0.000	**Jarque-Bera (JB):**	30.381
Skew:	1.249	**Prob(JB):**	2.53e-07
Kurtosis:	5.079	**Cond. No.**	7.12e+04

Warnings:
[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
[2] The condition number is large, 7.12e+04. This might indicate that there are
strong multicollinearity or other numerical problems.

According to our analysis, each additional mpg a domestic car has, the price increases on average 236 dollars per mile. If a car is foreign, the price increases by 1,072 on average and the with each additional mpg the average car price decreases by 37 dollars.